

Name: Harshita Singh
Class: D15C
Roll no. : 64
Subject: Machine Learning & Deep Learning

Experiment 7: Build an Artificial Neural Network (ANN) using Keras/TensorFlow

1. Dataset Source

- **Dataset Name:** Parkinson's Disease Data Set
 - **Source Repository:** UCI Machine Learning Repository
 - **Principal Investigator:** Max Little, University of Oxford
 - **Official Link:** <https://archive.ics.uci.edu/dataset/174/parkinsons>
 - **Data File Link:**
<https://archive.ics.uci.edu/ml/machine-learning-databases/parkinsons/parkinsons.data>
-

2. Dataset Description

Overview

This dataset is composed of a range of biomedical voice measurements from 31 individuals, 23 of whom have been diagnosed with Parkinson's Disease (PD). The primary objective of the dataset is to differentiate healthy individuals from those with PD using non-invasive vocal biomarkers.

Size and Structure

- **Total Instances:** 195 voice recordings.
- **Attributes:** 24 columns (23 vocal features + 1 target variable).
- **Format:** Multivariate, real-valued attributes in CSV (ASCII) format.

Target Variable

- **Status:** A binary value where **0** represents a healthy individual and **1** represents a patient with Parkinson's Disease.

Feature Categories

The dataset includes several highly technical measurements extracted from voice signals:

- **Fundamental Frequency:** MDVP:Fo(Hz), MDVP:Fhi(Hz), and MDVP:Flo(Hz) (Average, Maximum, and Minimum vocal frequencies).

- **Jitter Measures:** Five measures of variation in fundamental frequency (e.g., MDVP:Jitter(%), Jitter:DDP), which indicate "shakiness" in the voice.
- **Shimmer Measures:** Six measures of variation in amplitude (e.g., MDVP:Shimmer, Shimmer:APQ5), representing fluctuations in volume.
- **Noise-to-Harmonics Ratios:** NHR and HNR, which measure the ratio of noise to tonal components in the speech.
- **Nonlinear Complexity Measures:** RPDE, D2, DFA, and PPE (Pitch Period Entropy), used to capture the chaotic nature of vocal fold vibration common in neurological decline.

3. Mathematical Formulation of the Algorithm

An Artificial Neural Network (ANN) is a supervised learning model inspired by the biological structure of the human brain. The primary objective of an ANN is to approximate a complex mapping function between input features and output labels by minimizing prediction error.

The mathematical formulation can be understood by building from a single neuron up to the full learning cycle.

A. The Foundation: The Artificial Neuron

The fundamental building block of an ANN is the artificial neuron (or perceptron). Given an input dataset where an instance has n features, $X = \{x_1, x_2, \dots, x_n\}$, each neuron performs two distinct mathematical operations:

1. Linear Combination: The neuron calculates a weighted sum of its inputs and adds a bias term. This creates a linear output score, z :

$$z = \sum_{i=1}^n (w_i \cdot x_i) + b$$

- x_i : The input features (e.g., vocal frequency measurements).
- w_i : The weights assigned to each input (representing the importance of that feature).
- b : The bias term (an offset that shifts the activation threshold).

2. Activation (Non-Linearity):

To allow the network to model complex, non-linear relationships, the linear score z is passed through a mathematical filter called an activation function, resulting in the neuron's final output, a

$$a = f(z)$$

B. Activation Functions Used in this Model

Different layers of the network require different activation functions based on their purpose:

- **ReLU (Rectified Linear Unit):** Used in the **hidden layers**. It introduces non-linearity by simply outputting the input directly if it is positive, and outputting zero if it is negative. This helps the model learn faster and prevents the vanishing gradient problem.
 $f(z) = \max(0, z)$
- **Sigmoid Function:** Used in the **final output layer**. Because this experiment is a binary classification task (Healthy vs. Parkinson's), the Sigmoid function is used to squash the final network output into a probability between \$0\$ and \$1\$.

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

C. Network Architecture (Forward Propagation)

When multiple neurons are connected in layers, the output of one layer becomes the input of the next. For a specific layer l , the matrix formulation is:

$$Z^{[l]} = W^{[l]} A^{[l-1]} + b^{[l]}$$

$$A^{[l]} = f(Z^{[l]})$$

Where $W^{[l]}$ is the weight matrix for layer l , and $A^{[l-1]}$ is the activation output from the previous layer.

D. The Loss Function (Measuring Error)

Once the network makes a prediction ($\hat{y}$) for a patient, it must measure how far off that prediction is from the actual truth (y). For binary classification, we use the **Binary Cross-Entropy Loss** function. It heavily penalizes the model if it is highly confident in a wrong answer:

$$L(\hat{y}, y) = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

- N : Total number of samples.
- y_i : The actual true label (0 or 1).
- $\hat{y}_i$: The predicted probability that the label is 1.

E. Optimization and Backpropagation (The Learning Process)

To minimize the Loss (L), the network must learn by adjusting its weights. It does this by calculating the gradient (the derivative of the loss with respect to the weights) using the Chain Rule of calculus—a process called Backpropagation.

The weights are then updated in the opposite direction of the gradient using an optimization algorithm (such as Stochastic Gradient Descent or **Adam**):

$$W_{new} = W_{old} - \eta \cdot \frac{\partial L}{\partial W}$$

Where η (eta) represents the **learning rate**, controlling how large of a step the model takes when updating its weights.

4. Algorithm Limitations


While ANNs are powerful, they possess specific limitations that were evident during this experiment:

- **Data Requirement:** ANNs typically require large volumes of data to generalize. On small datasets (like this 195-row Parkinson's set), they are highly prone to **overfitting**, necessitating heavy regularization like Dropout.
- **Black-Box Nature:** Unlike Decision Trees or Logistic Regression, ANNs are not easily interpretable. It is difficult to explain exactly *which* vocal frequency feature triggered a "Parkinson's Positive" result.
- **Computational Expense:** Training ANNs requires more memory and processing power (often requiring GPUs for larger models) compared to traditional algorithms like SVM.
- **Hyperparameter Sensitivity:** The model's success is heavily dependent on the "Trial and Error" of selecting the right number of layers, neurons, and learning rates.
- **Lack of Spatial/Sequential Logic:** Standard feed-forward ANNs treat all features as independent and do not inherently understand the time-series nature of audio signals unless specialized architectures (like RNNs) are used.

5. Methodology / Workflow

The experiment followed a systematic end-to-end machine learning pipeline:

Building an ANN with Keras/TensorFlow: A Step-by-Step Workflow



Made with Napkin

1. **Data Ingestion:** The `parkinsons.data` file was imported using Pandas. The "name" attribute was removed to prevent the model from learning patient IDs rather than biomedical features.
2. **Exploratory Data Analysis (EDA):** Performed correlation analysis to identify which vocal features (like PPE and Spread1) had the strongest linear relationship with the disease status.
3. **Data Preprocessing:**
 - **Stratified Splitting:** Divided data into 80% training and 20% testing while maintaining the class ratio.
 - **Standardization:** Applied **StandardScaler** to ensure all features had a mean of 0 and a standard deviation of 1. This prevents features with large ranges (like Frequency in Hz) from dominating features with small ranges (like Jitter percentages).
4. **Architectural Design:** Built a Sequential model in Keras consisting of two hidden layers (32 and 16 neurons) and a single-neuron Sigmoid output layer.
5. **Regularization Implementation:** Integrated **Dropout layers** (30% and 20%) to prevent co-adaptation of neurons and **Early Stopping** to halt training when validation loss stopped improving.
6. **Model Training:** Compiled with the `adam` optimizer and `binary_crossentropy` loss, training for a maximum of 150 epochs with a batch size of 16.
7. **Evaluation & Tuning:** Final assessment was conducted using a Confusion Matrix and Classification Report (Precision, Recall, F1-Score). Post-experiment, **Keras Tuner** was used to mathematically verify the optimal hyperparameters.

Code:

```
# =====  
# 1. IMPORT NECESSARY LIBRARIES  
# =====  
import pandas as pd
```

```

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Scikit-Learn for preprocessing and evaluation
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score

# TensorFlow/Keras for the Artificial Neural Network
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping

# =====
# 2. LOAD AND EXPLORE THE DATA
# =====
print("Loading dataset...")
url =
"https://archive.ics.uci.edu/ml/machine-learning-databases/parkinsons/parkinsons.data"
df = pd.read_csv(url)

# Visualization 1: Class Distribution
plt.figure(figsize=(6, 4))
sns.countplot(x='status', data=df, palette='viridis')
plt.title('Distribution of Parkinson\'s Status (0: Healthy, 1: PD)')
plt.xlabel('Status')
plt.ylabel('Count')
plt.show()

# Visualization 2: Correlation Heatmap of top features
plt.figure(figsize=(10, 8))
# Selecting a subset of features for a cleaner heatmap
cols_to_plot = ['MDVP:Fo(Hz)', 'MDVP:Fhi(Hz)', 'MDVP:Flo(Hz)',
'MDVP:Jitter(%)', 'MDVP:Shimmer', 'HNR', 'status']
sns.heatmap(df[cols_to_plot].corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap of Key Vocal Features')
plt.show()

# =====
# 3. DATA PREPROCESSING
# =====
# Drop the 'name' column as it's just an ID and holds no predictive value
X = df.drop(['name', 'status'], axis=1)
y = df['status']

```

```

# Split into Training (80%) and Testing (20%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42, stratify=y)

# Feature Scaling: Neural Networks are highly sensitive to unscaled data.
# We fit the scaler ONLY on the training data to prevent data leakage.
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# =====
# 4. BUILD THE ARTIFICIAL NEURAL NETWORK
# =====
print("\nBuilding the ANN Model...")
model = Sequential()

# Input Layer & 1st Hidden Layer
# 32 Neurons, ReLU activation function
model.add(Dense(32, activation='relu',
input_shape=(X_train_scaled.shape[1],)))
# Dropout layer: Randomly drops 30% of neurons to prevent overfitting
model.add(Dropout(0.3))

# 2nd Hidden Layer
model.add(Dense(16, activation='relu'))
model.add(Dropout(0.2))

# Output Layer
# 1 Neuron, Sigmoid activation (outputs a probability between 0 and 1 for
binary classification)
model.add(Dense(1, activation='sigmoid'))

# Compile the Model
# Adam optimizer is highly efficient. Binary Crossentropy is the standard loss
function for Yes/No tasks.
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

# Print the architecture
model.summary()

# =====
# 5. TRAIN THE MODEL
# =====
# Early stopping will monitor validation loss and stop training if it doesn't
improve for 15 epochs
early_stop = EarlyStopping(monitor='val_loss', patience=15,
restore_best_weights=True)

```

```

print("\nTraining the Model...")
history = model.fit(
    X_train_scaled,
    y_train,
    epochs=150,          # Maximum number of passes
    batch_size=16,       # Number of samples per gradient update
    validation_split=0.2, # Use 20% of training data to evaluate during
training
    callbacks=[early_stop],
    verbose=1            # Set to 0 if you don't want the epoch loading bars
to print
)

# =====
# 6. VISUALIZE TRAINING PERFORMANCE
# =====
# Create a figure with two subplots: Loss and Accuracy
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

# Plot Training vs Validation Loss
ax1.plot(history.history['loss'], label='Training Loss')
ax1.plot(history.history['val_loss'], label='Validation Loss')
ax1.set_title('Model Loss Over Epochs')
ax1.set_xlabel('Epoch')
ax1.set_ylabel('Loss (Binary Crossentropy)')
ax1.legend()

# Plot Training vs Validation Accuracy
ax2.plot(history.history['accuracy'], label='Training Accuracy')
ax2.plot(history.history['val_accuracy'], label='Validation Accuracy')
ax2.set_title('Model Accuracy Over Epochs')
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Accuracy')
ax2.legend()

plt.show()

# =====
# 7. EVALUATE ON TEST DATA
# =====
print("\nEvaluating Model on Unseen Test Data...")
# Predict probabilities for the test set
y_pred_probs = model.predict(X_test_scaled)
# Convert probabilities to binary class labels (0 or 1) using a 0.5 threshold
y_pred = (y_pred_probs > 0.5).astype(int)

# Print metrics
print(f"\nTest Accuracy: {accuracy_score(y_test, y_pred) * 100:.2f}%\n")
print("--- Classification Report ---")

```



```

print(classification_report(y_test, y_pred))

# Visualization 3: Confusion Matrix
plt.figure(figsize=(6, 5))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Healthy (0)', 'Parkinson\'s (1)'],
            yticklabels=['Healthy (0)', 'Parkinson\'s (1)'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

# =====
# 8. HYPERPARAMETER TUNING WITH KERAS TUNER
# =====
import keras_tuner as kt

#. Define a "Model Builder" function
def model_builder(hp):
    model = tf.keras.Sequential()

    # Tune the number of units in the first Dense layer
    # Choose an optimal value between 16-64
    hp_units = hp.Int('units', min_value=16, max_value=64, step=16)
    model.add(tf.keras.layers.Dense(units=hp_units, activation='relu',
input_shape=(X_train_scaled.shape[1],)))

    # Tune the dropout rate
    hp_dropout = hp.Float('dropout', min_value=0.1, max_value=0.5, step=0.1)
    model.add(tf.keras.layers.Dropout(rate=hp_dropout))

    model.add(tf.keras.layers.Dense(16, activation='relu'))
    model.add(tf.keras.layers.Dense(1, activation='sigmoid'))

    # Tune the learning rate for the optimizer
    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=hp_learning_rate),
            loss='binary_crossentropy',
            metrics=['accuracy'])

    return model

# Initialize the Tuner (RandomSearch)
tuner = kt.RandomSearch(
    model_builder,

```

```

    objective='val_accuracy',
    max_trials=5,          # It will try 5 different combinations
    directory='my_dir',
    project_name='intro_to_kt'
)

# 3. Run the Search
print("\n--- Starting Automated Hyperparameter Tuning ---")
tuner.search(X_train_scaled, y_train, epochs=50, validation_split=0.2)

# 4. Get the best hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials=1)[0]

print(f"""
The hyperparameter search is complete.
- The optimal number of units in the first layer is: {best_hps.get('units')}
- The optimal dropout rate is: {best_hps.get('dropout')}
- The optimal learning rate is: {best_hps.get('learning_rate')}
""")

```

OUTPUT:



Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 32)	736
dropout_2 (Dropout)	(None, 32)	0
dense_4 (Dense)	(None, 16)	528
dropout_3 (Dropout)	(None, 16)	0
dense_5 (Dense)	(None, 1)	17

Total params: 1,281 (5.00 KB)

Trainable params: 1,281 (5.00 KB)

Non-trainable params: 0 (0.00 B)




```
Evaluating Model on Unseen Test Data...
2/2 ————— 0s 50ms/step

Test Accuracy: 94.87%

--- Classification Report ---
              precision    recall  f1-score   support

      0       1.00      0.80      0.89         10
      1       0.94      1.00      0.97         29

   accuracy              0.95         39
  macro avg              0.97      0.90      0.93         39
 weighted avg              0.95      0.95      0.95         39
```



6. Performance Analysis

The Artificial Neural Network (ANN) was evaluated on a reserved testing set comprising 20% of the dataset (39 instances). The model achieved an outstanding overall **Test Accuracy of 94.87%**, demonstrating excellent generalization capabilities on unseen data. The performance is further analyzed across three key dimensions:

A. Training Dynamics and Generalization

The training history visualizations ("Loss Over Epochs" and "Accuracy Over Epochs") indicate a highly stable learning process. Both the training and validation loss converged smoothly without significant divergence, stabilizing around the 120th epoch. The presence of Dropout layers (rates of 0.3 and 0.2) effectively mitigated overfitting. Notably, the validation metrics tracked closely with the training metrics, proving that the network successfully learned the underlying feature representations of Parkinson's disease rather than merely memorizing the training samples.

B. Confusion Matrix Interpretation

In medical diagnostic models, evaluating the exact types of misclassifications is critical. The confusion matrix yields the following breakdown:

- **True Positives (29):** The model correctly identified all 29 instances of Parkinson's disease.
- **True Negatives (8):** The model correctly identified 8 healthy patients.
- **False Positives (2):** The model incorrectly classified 2 healthy individuals as having Parkinson's (Type I Error).
- **False Negatives (0):** The model did not misclassify a single Parkinson's patient as healthy (Type II Error).

C. Evaluation Metrics (Classification Report)

The classification report confirms the model's high reliability, particularly in detecting the disease:

- **Recall (Sensitivity) for Class 1:** The model achieved a perfect **1.00 (100%) recall** for Parkinson's detection. This is the most vital metric for a healthcare screening tool, as it ensures zero missed diagnoses (0 False Negatives).
- **Precision for Class 1:** The precision stands at **0.94 (94%)**, indicating that when the model predicts Parkinson's, it is correct 94% of the time, with only a minor trade-off of 2 false alarms.
- **F1-Score:** The harmonic mean of precision and recall for the positive class is **0.97**, showcasing an exceptionally well-balanced and highly accurate predictive model suitable for early biomedical voice screening.

7. Hyperparameter Tuning

Artificial Neural Networks are highly sensitive to hyperparameter choices, which directly affect training stability, accuracy, and generalization. Since manual selection of architecture can be biased, systematic tuning is necessary to discover the optimal mathematical configuration for a specific dataset. In this experiment, automated hyperparameter tuning was performed using the **Keras Tuner** library to explore the search space and select the most robust model based on validation performance.

The Tuned Hyperparameters:

- **Number of Hidden Neurons:** The capacity of the first hidden layer was tuned within a range of 16 to 64. While larger networks increase learning capacity, they risk overfitting on small datasets like the Parkinson's Data Set. The tuner identified **16 neurons** as the optimal balance for this specific feature set.
- **Dropout Rate:** To ensure the model did not simply memorize the 195 training instances, the dropout rate was tuned between 0.1 and 0.5. A rate of **0.3** was selected, providing the necessary regularization to ensure the model generalizes well to unseen patient data.
- **Learning Rate:** This controls the step size during gradient descent. Among the candidate values (0.01, 0.001, and 0.0001), the tuner selected **0.01**, allowing for efficient convergence given the relatively simple architecture.
- **Search Strategy:** A **Random Search** strategy was employed to sample the multidimensional search space. This approach is computationally efficient and effectively identifies high-performing regions of the hyperparameter space without the exhaustive cost of a full grid search.

Result Analysis of Tuning

The automated hyperparameter search was completed in **34 seconds** across 5 distinct trials. The analysis of the tuning process reveals the following insights:

```
Trial 5 Complete [00h 00m 05s]
val_accuracy: 0.96875

Best val_accuracy So Far: 0.96875
Total elapsed time: 00h 00m 34s

The hyperparameter search is complete.
- The optimal number of units in the first layer is: 16
- The optimal dropout rate is: 0.30000000000000004
- The optimal learning rate is: 0.01
```

1. **Validation Performance:** The best performing model achieved a **Validation Accuracy of 96.88%**. This is a slight improvement over the initial manual configuration, proving that reducing the complexity of the first layer (from 32 down to 16 units) actually improved the model's ability to generalize.

2. **Efficiency vs. Accuracy:** The tuner favored a **Learning Rate of 0.01**. On this small dataset, a higher learning rate allowed the model to reach an optimal loss state very quickly (as seen in the 5-second trial time) without overshooting the global minimum.
3. **Regularization Impact:** The selection of a **0.3 Dropout rate** confirms that the model required moderate "noise" during training to prevent it from becoming overly dependent on specific biomedical features, ensuring the 94-96% accuracy is sustainable on new data.
4. **Final Conclusion:** The tuning process confirmed that for the Parkinson's Disease Data Set, a **leaner architecture** (16 units) is superior to a dense one. By automating this selection, we have mathematically verified that the model is not just accurate by chance, but is configured optimally for the specific variance found in vocal frequency data.